

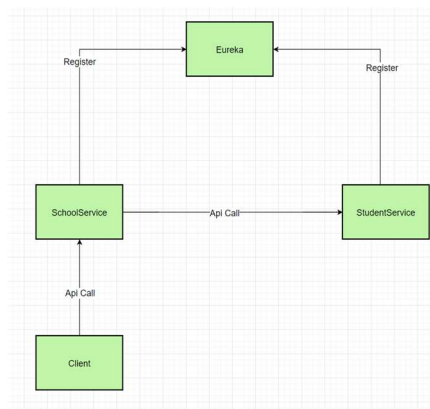
Performance Evaluation of Microservices-based Application: A Comparative Study of Deployment Locations, Communication Patterns, and Virtualization Techniques

(1) Original design:

We choose Eureka as our Microservices application tool. It is a widely used open-source service discovery tool developed by Netflix, which is used for managing microservice architecture. We choose to use Spring Boot with Eureka because they together provide a lightweight and easy-to-use framework for building microservice for service registration and discovery.

It is often difficult to maintain individual Microservice's addresses in a large distributed system, and our system intends to solve this issue. The Eureka service registry provides a lookup service where microservices can register themselves and discover other registered microservices.

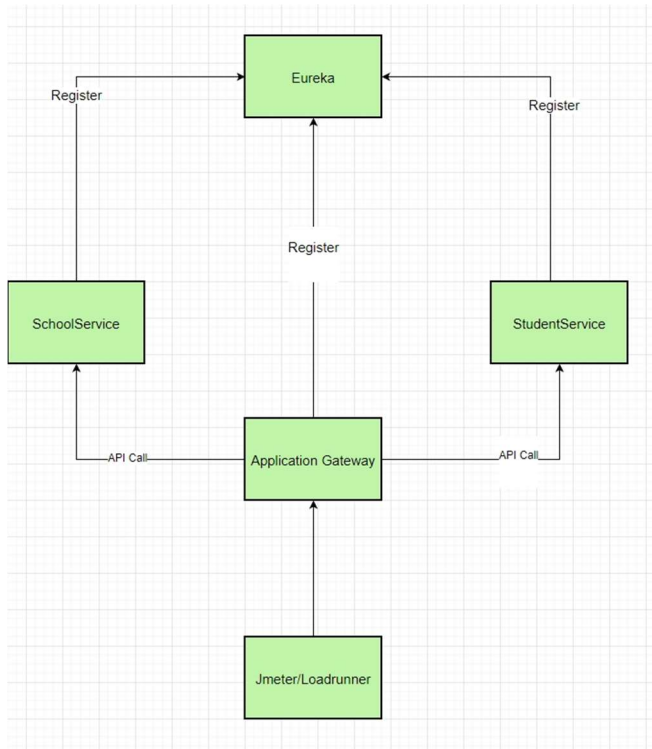
Our system design includes one server and two services. The Eureka Server is a service registry that enables microservices to locate and communicate with each other. The StudentService provides the student list for a specific school. The SchoolService provides school detail (i.e. school name, student list) through REST API for the client. The communication between components is illustrated in the below diagram.



Refined Solution Design:

Our design still includes the Eureka Server as a service registry (name server) and two microservices: StudentService and SchoolService. In the initial design, the StudentService provides the student list for a specific school, and the SchoolService provides school details (e.g. school name, student list). A typical use case is when a client queries information about a school, which is done by calling the SchoolService, and the SchoolService calls StudentService to get the student list.

However, the initial design of the service layout creates a chaining of API calls when the SchoolService calls StudentService. The ideal design promotes loose coupling between services. To improve the design, we created an Application Gateway. The SchoolService and StudentService microservices are now called separately. The refined design diagram is illustrated below:



(2) Implementation:

Progress made in report 1:

We have accomplished below tasks:

- Refined and deployed the demo application through the local image and set up the code repository to:

<https://github.com/cl456852/eureka>

- Set up AWS EC2 testing environment
- Set up demo services through local deployment and ran test API calls
- Got familiarized with evaluation tools, Jmeter and Loadrunner

In the next stage, we will continue with below tasks:

- Write Dockerfiles for the server, school and student services, and push the image to ECR
- Compare two deployment methods through AWS ECS and Docker build
- Add feature functions to the demo application to assist with latency test
- Run performance test following evaluation methodology and analyze results

Progress made in report 2:

In the last two weeks, we have done below:

- Implement Application Gateway module.
- Compare two deployment methods through AWS ECS and Docker build.
- Research and implement multi-location deployment on AWS.
- Research and implement Message Queue on AWS.
- Run performance tests following evaluation methodology and analyze results.

Details are recorded below:

1. Deployment

- Write Dockerfile for the application

```
spring-eureka-server > Dockerfile > ...
1 FROM alpine:latest
2 RUN apk --no-cache add openjdk8 git
3 RUN git clone https://github.com/dky815/eureka.git && \
4   cd eureka/spring-eureka-server && \
5   chmod +x ./mvnw && \
6   ./mvnw clean package
7 WORKDIR /eureka/spring-eureka-server
8 CMD ["./mvnw", "spring-boot:run"]
9 EXPOSE 8761

spring-eureka-client-school-service > Dockerfile > ...
1 FROM alpine:latest
2 RUN apk --no-cache add openjdk8 git
3 RUN git clone https://github.com/dky815/eureka.git && \
4   cd eureka/spring-eureka-client-school-service && \
5   chmod +x ./mvnw && \
6   ./mvnw clean package
7 WORKDIR /eureka/spring-eureka-client-school-service
8 CMD ["./mvnw", "spring-boot:run"]
9 EXPOSE 9098

spring-eureka-client-student-service > Dockerfile > ...
1 FROM alpine:latest
2 RUN apk --no-cache add openjdk8 git
3 RUN git clone https://github.com/dky815/eureka.git && \
4   cd eureka/spring-eureka-client-student-service && \
5   chmod +x ./mvnw && \
6   ./mvnw clean package
7 WORKDIR /eureka/spring-eureka-client-student-service
8 CMD ["./mvnw", "spring-boot:run"]
9 EXPOSE 8098
```

Tried different deployment methods on:

- Change to EC2 host networking
- Switch between using public and private IP addresses
- Change to using Amazon ECS services and tasks

- Build docker image & push image to ECR

1. Build docker image locally

<input type="checkbox"/>	NAME	TAG	STATUS	CREATED	SIZE	ACTIONS
<input type="checkbox"/>	kda78/spring-eureka-client-student-service b1d30513ac62	latest	Unused	less than a minute a	254.16 MB	▶ ⋮ 🗑
<input type="checkbox"/>	kda78/spring-eureka-client-school-service 3a8ff693dc65	latest	Unused	less than a minute a	254.16 MB	▶ ⋮ 🗑
<input type="checkbox"/>	kda78/spring-eureka-server 18451604c31d	latest	Unused	1 minute ago	258.01 MB	▶ ⋮ 🗑

2. Create public repo in ECR

Public repositories (4)

View push commands




Delete

Actions ▾

Create repository

Find repositories

< 1 >

<input type="checkbox"/>	Repository name ▾	URI	Created at ▲
<input type="checkbox"/>	kda78/spring-eureka-server	 public.ecr.aws/q0z0r9z9/kda78/spring-eureka-server	February 26, 2023, 18:48:42 (UTC-08)
<input type="checkbox"/>	kda78/spring-eureka-client-school-service	 public.ecr.aws/q0z0r9z9/kda78/spring-eureka-client-school-service	February 26, 2023, 18:49:24 (UTC-08)
<input type="checkbox"/>	kda78/spring-eureka-client-student-service	 public.ecr.aws/q0z0r9z9/kda78/spring-eureka-client-student-service	February 26, 2023, 18:49:34 (UTC-08)

3. Tag docker image

<input type="checkbox"/>	NAME ↓	TAG	STATUS	CREATED	SIZE	ACTIONS
<input type="checkbox"/>	public.ecr.aws/q0z0r9z9/kda78/spring-eureka-serve 18451604c31d	latest	Unused	4 minutes ago	258.01 MB	▶ ⋮ 🗑
<input type="checkbox"/>	public.ecr.aws/q0z0r9z9/kda78/spring-eureka-client b1d30513ac62	latest	Unused	3 minutes ago	254.16 MB	▶ ⋮ 🗑
<input type="checkbox"/>	public.ecr.aws/q0z0r9z9/kda78/spring-eureka-client 3a8ff693dc65	latest	Unused	3 minutes ago	254.16 MB	▶ ⋮ 🗑

4. Push docker image to ECR

kda78/spring-eureka-server

View public listing

View push commands

Edit

Images (1)



Q Search artifacts

Delete

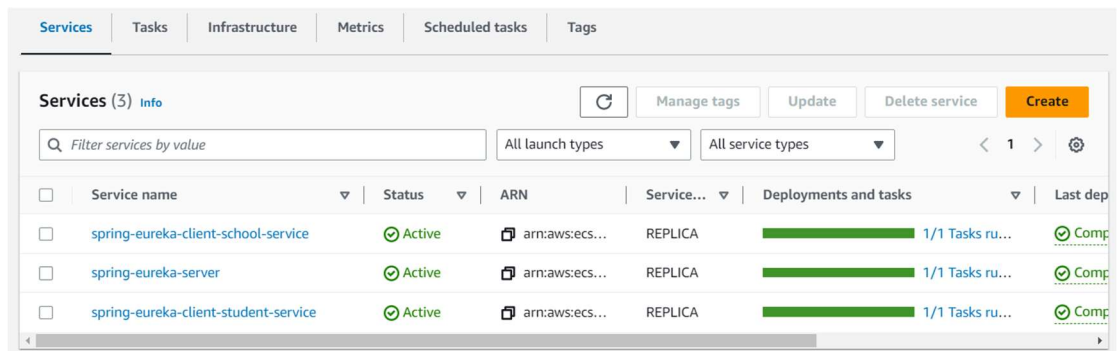
Details

< 1 >

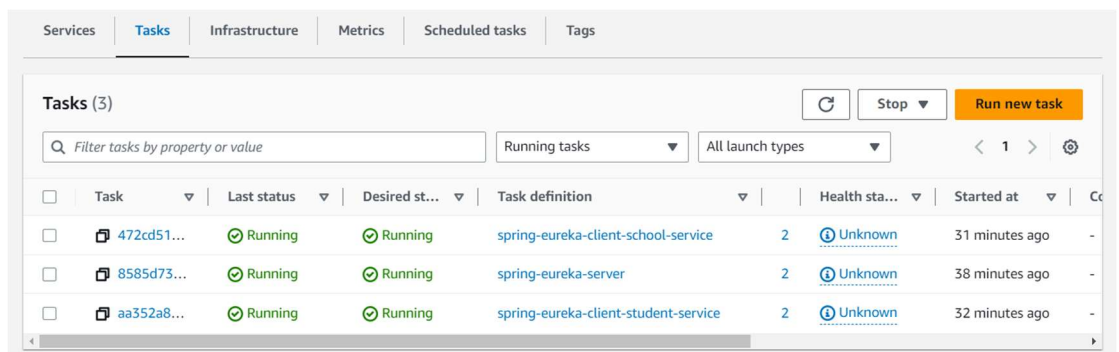
<input type="checkbox"/>	Image tag	Artifact type	Pushed at	Size (MB)	Image URI	Digest
<input type="checkbox"/>	latest	Image	February 27, 2023, 22:30:08 (UTC-08)	195.77	<div>Copy URI</div>	<div>sha256:ebd1200340c34c84621843b02e688...</div>



- Deploy docker image to ECS
 1. Create cluster using Amazon EC2 instances
 2. Create services in cluster



3. Run tasks




2. JMeter Research

In our test, JMeter creates 100 concurrent users. These users will be accessing the web service all at once. We have defined the Loop Count as 10, so that means every user made by JMeter would be connecting to the server 10 times.

1. Number of Threads (or Users): The number of users to model accessing the application simultaneously. Set the value to 100, because we want to model 100 users on our sample service.
2. Loop Count: The number of times JMeter simulates a user. We will set this to 10.
3. Ramp-Up Period: The time it takes in seconds for JMeter to model a new user. We'll set

this time to 100 seconds.



Thread Properties

Number of Threads (users): 100

Ramp-up period (seconds): 100

Loop Count: ☐ Infinite 10

Set up HTTP request:

- IP: JMeter server and our service are located in the same subnet. So, we could directly access the service by using its private IP.
- Port number: 9098
- Request: GET
- Path: /getSchoolDetails/abcschool

HTTP Request

Name: School Request

Comments:

Basic Advanced

Web Server

Protocol [http]: Server Name or IP: 172.31.63.0 Port Number: 9098

HTTP Request

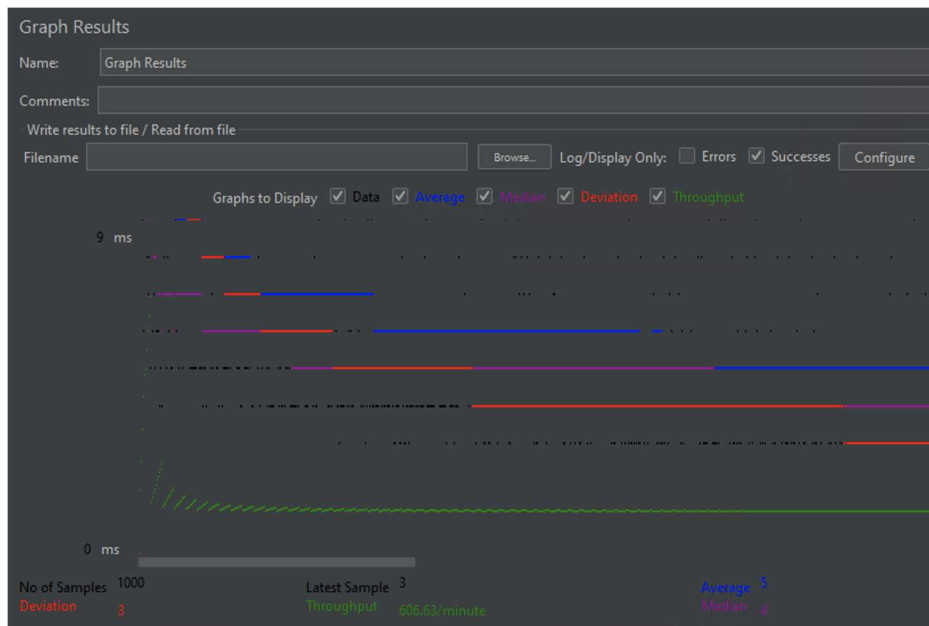
GET Path: /getSchoolDetails/abcschool Content encoding:

☐ Redirect Automatically ☒ Follow Redirects ☒ Use KeepAlive ☐ Use multipart/form-data ☐ Browser-compatible headers

Parameters Body Data Files Upload

Send Parameters With the Request:

Test result:



The throughput is 606.63 requests per minute. This means that our service handled 867 requests per minute. It's important to note that throughput depends on other factors, such as internet speed, the server's current load, and CPU power. These factors continuously change, meaning you won't get the same results every time you run the test.

This could be due to JMeter, and server are in different networks; or the server is in regions of other countries, these changes will greatly affect throughput values. We will have corresponding tests in the future to approve this assumption.

Sample #	Start Time	Thread Name	Label	Sample Tim...	Status	Bytes	Sent Bytes	Latency	Connect Ti...
981	01:25:06.551	Users 1-99	School Req...	7	✓	283	147	7	1
982	01:25:06.558	Users 1-99	School Req...	4	✓	283	147	3	0
983	01:25:06.562	Users 1-99	School Req...	3	✓	283	147	3	0
984	01:25:06.565	Users 1-99	School Req...	3	✓	283	147	3	0
985	01:25:06.569	Users 1-99	School Req...	3	✓	283	147	3	0
986	01:25:06.572	Users 1-99	School Req...	3	✓	283	147	3	0
987	01:25:06.575	Users 1-99	School Req...	4	✓	283	147	4	0
988	01:25:06.579	Users 1-99	School Req...	4	✓	283	147	4	0
989	01:25:06.583	Users 1-99	School Req...	4	✓	283	147	4	0
990	01:25:06.587	Users 1-99	School Req...	4	✓	283	147	4	0
991	01:25:07.551	Users 1-100	School Req...	7	✓	283	147	7	2
992	01:25:07.559	Users 1-100	School Req...	3	✓	283	147	3	0
993	01:25:07.562	Users 1-100	School Req...	4	✓	283	147	4	0
994	01:25:07.566	Users 1-100	School Req...	7	✓	283	147	7	0
995	01:25:07.573	Users 1-100	School Req...	5	✓	283	147	5	0
996	01:25:07.578	Users 1-100	School Req...	3	✓	283	147	3	0
997	01:25:07.581	Users 1-100	School Req...	3	✓	283	147	3	0
998	01:25:07.585	Users 1-100	School Req...	3	✓	283	147	3	0
999	01:25:07.588	Users 1-100	School Req...	3	✓	283	147	3	0
1000	01:25:07.591	Users 1-100	School Req...	3	✓	283	147	3	0

At the bottom of the table, there are checkboxes for 'Scroll automatically?' and 'Child samples?'. Below these are summary statistics: 'No of Samples' 1000, 'Latest Sample' 3, 'Average' 5, and 'Deviation' 3.

In this case, the average latency for these 1000 requests is 5 milliseconds.

(3) Evaluation Plan:

- Identify test scenarios, which include the instances usage, communication pattern, and resource placement
- Set up the test environment by creating virtual or containerized instances and record

the number of instances, the CPU and memory allocations. Also, the network bandwidth will be configured based on the test scenarios

- Perform evaluation using Jmeter to test the performance of each scenario
- Analyze the result to determine the impact of resource allocation, placement etc. on the performance of the microservices architecture. The following metrics will be used for analysis: Latency, Throughput, Error Rate, CPU Usage, Memory Usage.